

Imarticus Capstone Project By Saurabh Gupta

Importing Data and Libraries

In [1]:

```
#Importing Data and Libraries
import pandas as pd
import numpy as np
lc=pd.read_csv("S:\XYZCorp_LendingData(1).txt",sep = '\t',na_values = 'NaN',low_memory = False)
```

In [5]:

```
print(lc.head(4))
print(lc.shape)
y=lc['default_ind']
```

```
      id  member_id  loan_amnt  funded_amnt  funded_amnt_inv      term \
0  1077501    1296599    5000.0    5000.0         4975.0    36 months
1  1077430    1314167    2500.0    2500.0         2500.0    60 months
2  1077175    1313524    2400.0    2400.0         2400.0    36 months
3  1076863    1277178   10000.0   10000.0        10000.0    36 months

      int_rate  installment  grade  sub_grade  ...  il_util  open_rv_12m  open_rv_24m  \
0      10.65         162.87      B      B2      ...      NaN          NaN          NaN
1      15.27          59.83      C      C4      ...      NaN          NaN          NaN
2      15.96          84.33      C      C5      ...      NaN          NaN          NaN
3      13.49         339.31      C      C1      ...      NaN          NaN          NaN

      max_bal_bc  all_util  total_rev_hi_lim  inq_fi  total_cu_tl  inq_last_12m  \
0           NaN         NaN                NaN    NaN          NaN          NaN
1           NaN         NaN                NaN    NaN          NaN          NaN
2           NaN         NaN                NaN    NaN          NaN          NaN
3           NaN         NaN                NaN    NaN          NaN          NaN

      default_ind
0              0
1              1
2              0
3              0

[4 rows x 73 columns]
(855969, 73)
```

Splitting Data into train and test

In [6]:

```
import pandas as pd
from sklearn.model_selection import train_test_split
X_train, X_test, y_train, y_test = train_test_split(lc, y, test_size=0.30)
print( X_train.shape, y_train.shape)
print( X_test.shape, y_test.shape)
```

```
(599178, 73) (599178,)
(256791, 73) (256791,)
```

Exploratory Data Analysis

In [7]:

```
#Removing Columns
Range_columns = [i for i in range(59, 70)]
```

```
Range_columns = [1, 2, 3, 4, 5, 6, 7, 8, 9, 10, 11, 12, 13, 14, 15, 16, 17, 18, 19, 20, 21, 22, 23, 24, 25, 26, 27, 28, 29, 30, 31, 32, 33, 34, 35, 36, 37, 38, 39, 40, 41, 42, 43, 44, 45, 46, 47, 48, 49, 50, 51, 52, 53, 54, 55, 56, 57, 58, 59, 60, 61, 62, 63, 64, 65, 66, 67, 68, 69, 70, 71, 72, 73, 74, 75, 76, 77, 78, 79, 80, 81, 82, 83, 84, 85, 86, 87, 88, 89, 90, 91, 92, 93, 94, 95, 96, 97, 98, 99, 100]
Train_data_1 = X_train.drop(X_train.columns[Range_columns],axis = 1)
Test_data_1 = X_test.drop(X_test.columns[Range_columns],axis =1)
```

In [8]:

```
print("the shape of train data is :",Train_data_1.shape)
print("the shape of train data is :",Test_data_1.shape)
```

```
the shape of train data is : (599178, 62)
the shape of train data is : (256791, 62)
```

In [10]:

```
del_columns = ['total_cu_tl','inq_last_12m','annual_inc_joint','dti_joint',
               'verification_status_joint']
Train_data_1 = Train_data_1.drop(labels = del_columns,axis = 1)
Test_data_1 = Test_data_1.drop(labels = del_columns,axis = 1)
```

In [11]:

```
drop_more_col = ['open_il_6m']
Train_data_2 = Train_data_1.drop(labels = drop_more_col,axis = 1)
Test_data_2 = Test_data_1.drop(labels = drop_more_col,axis = 1)
```

In [12]:

```
print("the shape of train data;",Train_data_2.shape)
print("the shape of test data;",Test_data_2.shape)
```

```
the shape of train data; (599178, 56)
the shape of test data; (256791, 56)
```

In [13]:

```
mean_train = Train_data_2[['mths_since_last_delinq','mths_since_last_record','revol_util',
                           'collections_12_mths_ex_med','mths_since_last_major_derog','tot_coll_amt','tot_cur_bal']].mean()
mean_test = Test_data_2[['mths_since_last_delinq','mths_since_last_record','revol_util',
                          'collections_12_mths_ex_med','mths_since_last_major_derog','tot_coll_amt','tot_cur_bal']].mean()
```

In [14]:

```
Train_data_2[['mths_since_last_delinq','mths_since_last_record','revol_util','collections_12_mths_ex_med',
               'mths_since_last_major_derog','tot_coll_amt','tot_cur_bal']] = Train_data_2[['mths_since_last_delinq','mths_since_last_record','revol_util','collections_12_mths_ex_med',
               'mths_since_last_major_derog','tot_coll_amt','tot_cur_bal']].fillna(mean_train)
Test_data_2[['mths_since_last_delinq','mths_since_last_record','revol_util','collections_12_mths_ex_med',
              'mths_since_last_major_derog','tot_coll_amt','tot_cur_bal']] = Test_data_2[['mths_since_last_delinq','mths_since_last_record','revol_util','collections_12_mths_ex_med',
              'mths_since_last_major_derog','tot_coll_amt','tot_cur_bal']].fillna(mean_test)
```

In [15]:

```
pd.unique(Train_data_2[['term','grade','home_ownership']].values.ravel('K'))
pd.unique(Test_data_2[['term','grade','home_ownership']].values.ravel('K'))
```

Out[15]:

```
array([' 36 months', ' 60 months', 'B', 'C', 'D', 'A', 'E', 'G', 'F',
       'RENT', 'MORTGAGE', 'OWN', 'OTHER', 'ANY', 'NONE'], dtype=object)
```

In [16]:

```
Train_data_2['grade'] = Train_data_2['grade'].map({'A':7,'B':6,'C':5,'D':4,'E':3,'F':2,'G':1})
```

```

Train_data_2['home_ownership'] = Train_data_2['home_ownership'].map({'MORTGAGE':6, 'RENT':5, 'OWN':4, 'OTHER':3, 'NONE':2, 'ANY':1})

Train_data_2["emp_length"] = Train_data_2["emp_length"].replace({'years':'', 'year':'', '': '', '<':'', '\+':'', 'n/a':'0', 's':''},
                                                                regex = True)

##test
Test_data_2['grade'] = Test_data_2['grade'].map({'A':7, 'B':6, 'C':5, 'D':4, 'E':3, 'F':2, 'G':1})
Test_data_2['home_ownership'] = Test_data_2['home_ownership'].map({'MORTGAGE':6, 'RENT':5, 'OWN':4, 'OTHER':3, 'NONE':2, 'ANY':1})

Test_data_2["emp_length"] = Test_data_2["emp_length"].replace({'years':'', 'year':'', '': '', '<':'', '\+':'', 'n/a':'0', 's':''},
                                                                regex = True)

```

In [17]:

```

Train_data_2['term'] = Train_data_2['term'].replace({'months':'', '': ''}, regex = True)
Test_data_2['term'] = Test_data_2['term'].replace({'months':'', '': ''}, regex = True)

```

In [18]:

```
print(Train_data_2['emp_length'].value_counts(dropna = False))
```

```

10    197360
1      85714
2      53220
3      47411
5      37702
4      35428
7      30280
NaN    30075
8      29658
6      28981
9      23349
Name: emp_length, dtype: int64

```

In [19]:

```

Train_data_2['emp_length'] = pd.to_numeric(Train_data_2['emp_length'], errors = 'coerce')
Test_data_2['emp_length'] = pd.to_numeric(Test_data_2['emp_length'], errors = 'coerce')
mean = Train_data_2['emp_length'].mean()
mean_test = Test_data_2['emp_length'].mean()
Train_data_2['emp_length'] = Train_data_2['emp_length'].fillna(mean)
Test_data_2['emp_length'] = Test_data_2['emp_length'].fillna(mean_test)
Train_data_2['term'] = pd.to_numeric(Train_data_2['term'], errors = 'coerce')
Test_data_2['term'] = pd.to_numeric(Test_data_2['term'], errors = 'coerce')

```

Feature Selection

In [21]:

```

#feature Selection
features = ['loan_amnt', 'funded_amnt', 'funded_amnt_inv', 'term', 'int_rate', 'installment', 'grade', 'emp_length', 'home_ownership', 'annual_inc', 'dti', 'delinq_2yrs', 'inq_last_6mths', 'mths_since_last_delinq', 'mths_since_last_record', 'open_acc', 'pub_rec', 'revol_bal', 'revol_util', 'total_acc', 'out_prncp', 'out_prncp_inv', 'total_pymnt', 'total_pymnt_inv', 'total_rec_prncp', 'total_rec_int', 'total_rec_late_fee', 'recoveries', 'collection_recovery_fee', 'last_pymnt_amnt', 'collections_12_mths_ex_med', 'mths_since_last_major_derog', 'policy_code', 'acc_now_delinq', 'tot_coll_amt', 'tot_cur_bal']
train = Train_data_2[features]
test = Test_data_2[features]

```

In [22]:

```
test.head()
```

Out[22]:

	loan_amnt	funded_amnt	funded_amnt_inv	term	int_rate	installment	grade	emp_length	home_ownership	annual_inc	...	t
187717	7500.0	7500.0	7500.0	36	11.14	246.04	6	7.0	5	60000.0	...	
495273	13700.0	13700.0	13700.0	60	12.59	308.85	5	1.0	5	45000.0	...	
517597	12000.0	12000.0	12000.0	36	10.99	392.81	6	5.0	6	38000.0	...	
418953	16800.0	16800.0	16800.0	36	11.99	557.93	6	2.0	6	63000.0	...	
346762	21750.0	21750.0	21750.0	36	16.29	767.79	4	10.0	6	128200.0	...	

5 rows × 36 columns

Random Forest Classification

In [57]:

```
from sklearn.feature_selection import SelectFromModel
from sklearn.ensemble import RandomForestClassifier
select = SelectFromModel(RandomForestClassifier(n_estimators=100, random_state=42),
threshold='median')
```

In [58]:

```
select.fit(train,y_train)
```

Out[58]:

```
SelectFromModel(estimator=RandomForestClassifier(bootstrap=True, class_weight=None,
criterion='gini',
          max_depth=None, max_features='auto', max_leaf_nodes=None,
          min_impurity_decrease=0.0, min_impurity_split=None,
          min_samples_leaf=1, min_samples_split=2,
          min_weight_fraction_leaf=0.0, n_estimators=100, n_jobs=None,
          oob_score=False, random_state=42, verbose=0, warm_start=False),
          max_features=None, norm_order=1, prefit=False, threshold='median')
```

In [59]:

```
X_train_s = select.transform(train)
```

In [60]:

```
print('The shape of train is: ', train.shape)
print('The shape of X_train_s is ', X_train_s.shape)
```

```
The shape of train is: (599178, 36)
The shape of X_train_s is (599178, 18)
```

In [61]:

```
X_train_s =pd.DataFrame(X_train_s)
```

In [62]:

```
X_train_s.head()
```

Out[62]:

	0	1	2	3	4	5	6	7	8	9	10	11	12	13	14	15	1
0	8000.0	8000.0	8000.0	36.0	10.49	259.99	5182.60	5182.60	3635.20	3635.20	2817.40	817.80	0.0	0.0	0.0	259.99	0.

1	21000.0	21000.0	21000.0	60.0	20.99	568.01	20175.95	20175.95	2247.55	2247.55	824.05	1423.50	0.0	0.0	0.0	568.01	0.1
2	13800.0	13800.0	13800.0	36.0	8.90	438.20	4612.20	4612.20	10955.00	10955.00	9187.80	1767.20	0.0	0.0	0.0	438.20	0.1
3	10000.0	10000.0	10000.0	36.0	10.99	327.34	9764.24	9764.24	315.13	315.13	235.76	79.37	0.0	0.0	0.0	327.34	112.1
4	10000.0	10000.0	10000.0	60.0	12.29	223.92	9122.89	9122.89	1560.61	1560.61	877.11	683.50	0.0	0.0	0.0	223.92	0.1

In [63]:

```
train.head()
```

Out[63]:

	loan_amnt	funded_amnt	funded_amnt_inv	term	int_rate	installment	grade	emp_length	home_ownership	annual_inc	...	t
253650	8000.0	8000.0	8000.0	36	10.49	259.99	6	2.0	4	47900.0	...	
584010	21000.0	21000.0	21000.0	60	20.99	568.01	3	5.0	6	85000.0	...	
52427	13800.0	13800.0	13800.0	36	8.90	438.20	7	7.0	6	60000.0	...	
494114	10000.0	10000.0	10000.0	36	10.99	327.34	6	1.0	5	48000.0	...	
732309	10000.0	10000.0	10000.0	60	12.29	223.92	5	5.0	5	45000.0	...	

5 rows × 36 columns

In [64]:

```
import matplotlib.pyplot as plt
mask = select.get_support()
print(mask)
plt.matshow(mask.reshape(1,-1), cmap='gray_r')
plt.xlabel('Index of Features')
```

```
[ True  True  True  True  True  True False False False False False False
 False False False False False False False False  True  True  True  True
  True  True  True  True  True  True False False False False  True  True]
```

Out[64]:

Text(0.5, 0, 'Index of Features')



In [65]:

```
from sklearn.feature_selection import SelectPercentile
from sklearn.feature_selection import f_classif
Selector_f = SelectPercentile(f_classif, percentile=25)
Selector_f.fit(train,y_train)

for n,s in zip(train,Selector_f.scores_):
    print('F-score:', (s,n))
```

C:\Users\Saurabh Gupta\Anaconda3\lib\site-packages\sklearn\feature_selection\univariate_selection.py:114: UserWarning: Features [32] are constant.
UserWarning)
C:\Users\Saurabh Gupta\Anaconda3\lib\site-packages\sklearn\feature_selection\univariate_selection.py:115: RuntimeWarning: invalid value encountered in true_divide
f = msb / msw

```
F-score: (16.413181516820682, 'loan_amnt')
F-score: (23.12901430907968, 'funded_amnt')
F-score: (45.31975259795728, 'funded_amnt_inv')
F-score: (625.2272036418016, 'term')
F-score: (14812.306056010419, 'int_rate')
F-score: (9.300231133384154, 'installment')
```

```

F-score: (9339.229597699246, 'grade')
F-score: (105.4136765183886, 'emp_length')
F-score: (220.49116985892073, 'home_ownership')
F-score: (782.9154560984659, 'annual_inc')
F-score: (6.9034912247148075, 'dti')
F-score: (60.27347117621261, 'delinq_2yrs')
F-score: (3362.530712095092, 'inq_last_6mths')
F-score: (1.488617030312238, 'mths_since_last_delinq')
F-score: (306.3167076700082, 'mths_since_last_record')
F-score: (279.1596226299355, 'open_acc')
F-score: (250.32876354490753, 'pub_rec')
F-score: (266.3785707727901, 'revol_bal')
F-score: (1181.9835113566942, 'revol_util')
F-score: (257.0788506128066, 'total_acc')
F-score: (32450.020302320205, 'out_prncp')
F-score: (32449.544800335883, 'out_prncp_inv')
F-score: (931.0629294764716, 'total_pymnt')
F-score: (982.3834693029864, 'total_pymnt_inv')
F-score: (4969.582377032308, 'total_rec_prncp')
F-score: (1262.3139891354192, 'total_rec_int')
F-score: (12282.539587793462, 'total_rec_late_fee')
F-score: (173105.15754890683, 'recoveries')
F-score: (72929.56272535223, 'collection_recovery_fee')
F-score: (4624.280275711025, 'last_pymnt_amnt')
F-score: (78.40110717524166, 'collections_12_mths_ex_med')
F-score: (19.576990490565475, 'mths_since_last_major_derog')
F-score: (nan, 'policy_code')
F-score: (8.017872935748445, 'acc_now_delinq')
F-score: (41.14172910398002, 'tot_coll_amt')
F-score: (601.1655049113756, 'tot_cur_bal')

```

Logistic Regression

In [23]:

```

#logistic regression
from sklearn.linear_model import LogisticRegression

```

In [24]:

```

features = ['loan_amnt', 'funded_amnt', 'funded_amnt_inv', 'term', 'int_rate', 'installment', 'grade', 'ou
t_prncp',
            'out_prncp_inv', 'total_pymnt',
            'total_pymnt_inv', 'total_rec_prncp',
            'total_rec_int', 'total_rec_late_fee', 'recoveries', 'collection_recovery_fee',
            'last_pymnt_amnt', 'tot_cur_bal']

```

In [25]:

```

Final_train = Train_data_2[features]
Final_test = Test_data_2[features]

```

In [27]:

```
y_train.shape
```

Out[27]:

```
(599178,)
```

In [28]:

```
Final_train.shape
```

Out[28]:

```
(599178, 18)
```

In [29]:

In [31]:

```
X_train1,X_test1,y_train1,y_test1 = train_test_split(Final_train,y_train,random_state = 42,test_size = 0.3)
```

In [33]:

```
logreg = LogisticRegression()  
logreg.fit(X_train1,y_train1)
```

C:\Users\Saurabh Gupta\Anaconda3\lib\site-packages\sklearn\linear_model\logistic.py:433:
FutureWarning: Default solver will be changed to 'lbfgs' in 0.22. Specify a solver to silence this
warning.
FutureWarning)

Out[33]:

```
LogisticRegression(C=1.0, class_weight=None, dual=False, fit_intercept=True,  
                    intercept_scaling=1, max_iter=100, multi_class='warn',  
                    n_jobs=None, penalty='l2', random_state=None, solver='warn',  
                    tol=0.0001, verbose=0, warm_start=False)
```

In [34]:

```
y_pred = logreg.predict(X_test1)
```

In [36]:

```
logreg.score(X_test1,y_test1)
```

Out[36]:

0.9977024155234376

In [37]:

```
from sklearn.metrics import confusion_matrix  
from sklearn.metrics import classification_report
```

In [38]:

```
print("confusion matrix is: {}".format(confusion_matrix(y_test1,y_pred)))
```

```
confusion matrix is: [[169839      5]  
 [   408   9502]]
```

In [39]:

```
print(classification_report(y_test1,y_pred))
```

	precision	recall	f1-score	support
0	1.00	1.00	1.00	169844
1	1.00	0.96	0.98	9910
micro avg	1.00	1.00	1.00	179754
macro avg	1.00	0.98	0.99	179754
weighted avg	1.00	1.00	1.00	179754

In [40]:


```
#plotting the roc curve  
from sklearn.metrics import roc_curve  
import matplotlib.pyplot as plt  
from sklearn.metrics import roc_auc_score
```

In [41]:

```
y_pred_prob = logreg.predict_proba(X_test1)[: ,1]
fpr,tpr,threshold = roc_curve(y_test1,y_pred_prob)
plt.xlabel('Fpr')
plt.ylabel('Tpr')
plt.plot([0,1],[0,1], 'k--')
plt.plot(fpr,tpr,label = 'logistic regression')
```

Out[41]:

[<matplotlib.lines.Line2D at 0x29d87cf2ac8>]



In [42]:

```
roc_auc_score(y_test1,y_pred_prob)
```

Out[42]:

0.987686109822723

In [43]:

```
#predicting on the final test values
logreg.predict(Final_test)[0:5]
```

Out[43]:

array([0, 0, 0, 0, 0], dtype=int64)

In [44]:

```
Predict_finaltest = logreg.predict_proba(Final_test)[: ,1]
print("The probability of the coming 1 on the final_tes is :{}".format(Predict_finaltest,2))
```

The probability of the coming 1 on the final_tes is :[5.69948607e-07 5.98333912e-04 2.56170542e-04 ... 5.01892091e-05 2.07613474e-25 5.99070912e-13]

In [45]:

```
roc_auc_score(y_test,Predict_finaltest)
```

Out[45]:

0.9881555378324336

In [46]:


```
Predict_finaltest = logreg.predict_proba(Final_test)[: ,1]
roc_curve(y_test,Predict_finaltest)
plt.xlabel('Fpr')
plt.ylabel('Tpr')
plt.plot([0,1],[0,1], 'k--')
```



```
plt.plot([0,1],[0,1], 'k--')
plt.plot(fpr,tpr,label = 'logistic regression')
```

Out[46]:

[<matplotlib.lines.Line2D at 0x29d8193b6d8>]



Knn Classification

In [52]:

```
#knn classification
from sklearn.neighbors import KNeighborsClassifier
X_train1,X_test1,y_train1,y_test1 = train_test_split(Final_train,y_train,random_state = 42,test_size = 0.3)
knn = KNeighborsClassifier(n_neighbors=5)
knn.fit(X_train1,y_train1)
y_pred = knn.predict(X_test1)
knn.score(X_test1,y_test1)
```

Out[52]:


0.9957664363519032

In [53]:

```
knn_predict = knn.predict_proba(X_test1)[:,:1]
fpr,tpr,threshold = roc_curve(y_test1,knn_predict)
plt.xlabel('Fpr')
plt.ylabel('Tpr')
plt.plot([0,1],[0,1], 'k--')
plt.plot(fpr,tpr,label = 'knn')
```

Out[53]:

[<matplotlib.lines.Line2D at 0x29d87dff940>]



In [54]:

```
#predicting on the final test set using knn..  
y_pred = knn.predict(Final_test)  
y_pred
```

Out[54]:

```
array([0, 0, 0, ..., 0, 0, 0], dtype=int64)
```

In [55]:

```
pred = knn.predict_proba(Final_test)[:,-1]
```

In [56]:

```
roc_auc_score(y_test,pred)
```

Out[56]:

```
0.9781692320503372
```